

Crowd Sourcing FAQ Project Report

Samagama: AI-Powered Semantic FAQ and Crowdsourcing Platform

Team Members

Radhika Sharma (Team lead)

Nada Waris Gilani

Dhruvi Shah

Mansi Singh

Anushka Jaiswal

Abhishek Yadav

Piyush Yadav

Nisha Chendke

Mangali Bhaskar

Bhoomika Banjaran

Repository Name: faq

Repository Link: <https://github.com/ABHISHEK27Y/faq.git>

June 22, 2026

Contents

1	Executive Summary	1
2	Introduction	2
2.1	Background and Motivation	2
2.2	Proposed Solution: Samagama	2
2.3	System Requirements	3
3	System Design & Architecture	4
3.1	High-Level Architecture	4
3.2	Entity Relationship Diagram (ERD)	5
3.3	Use-Case Diagram	5
3.3.1	Detailed Use Case Specification: RAG Auto-Answer Generation	6
3.4	Machine Learning Inference Pipeline	8
3.5	API Route & Middleware Design	8
3.6	End-to-End System Pipeline Flow	8
4	Data Flow Diagrams (DFD)	10
4.1	Level 0 DFD (Context Diagram)	10
4.2	Level 1 DFD	10
5	Implementation	12
5.1	Tech Stack & Justification	12
5.2	API Design & Endpoints	13
5.3	Module & Feature Breakdown	13
6	Feature Spotlight	15
6.1	Semantic Vector Search (Intent Engine)	15
6.1.1	Theoretical Purpose	15
6.1.2	Architecture and Execution Algorithm	15
6.1.3	Real-World Impact	16
6.2	Deep Auto-Answer Pipeline (Yaksha Bot)	17
6.2.1	Theoretical Purpose	17
6.2.2	Sequence Diagram & RAG Flow	17
6.2.3	Architecture and RAG Algorithm	17
6.2.4	Real-World Impact	18
6.3	Synchronous AI Toxicity Moderation	19

6.3.1	Theoretical Purpose	19
6.3.2	Architecture and Execution Algorithm	19
6.3.3	Real-World Impact	20
6.4	Expert Promotion & Reputation Engine	21
6.4.1	Theoretical Purpose	21
6.4.2	Architecture and Execution Algorithm	21
6.4.3	Real-World Impact	21
6.5	Live Socket.io Real-Time Telemetry	22
6.5.1	Theoretical Purpose	22
6.5.2	Architecture and Execution Algorithm	22
6.5.3	Real-World Impact	22
7	Challenges & System Limitations	23
7.1	Challenge 1: AI Latency vs. UX Blocking	23
7.2	Challenge 2: Vector Dimensionality Mismatch	23
7.3	Limitation: In-Memory Array Iteration	23
8	Future Architectural Enhancements	24
9	Conclusion	25

1 Executive Summary

The Samagama FAQ & Crowdsourcing Platform represents a paradigm shift in institutional knowledge management. Modern enterprises, educational institutions, and public service domains consistently grapple with the challenge of disseminating accurate policy information. Traditional Frequently Asked Question (FAQ) systems, while ubiquitous, suffer from an inherent lexical matching limitation. When end-users query these systems using synonyms, abstract phrasing, or colloquial language, legacy keyword-matching algorithms invariably fail. This results in the "No Results Found" phenomenon, leading to user frustration and the exponential inflation of administrative support tickets.

To definitively solve this, Samagama discards traditional lexical search in favor of a state-of-the-art **Semantic Vector Search** (Intent Engine) powered by Google's Generative AI `text-embedding-004` model. This allows the system to comprehend the multidimensional geometric meaning behind user queries, ensuring high-fidelity retrieval regardless of linguistic variations.

Furthermore, to address the delayed resolution times inherent in decentralized community Q&A forums, the platform introduces a **Deep Auto-Answer Pipeline**. Utilizing the Retrieval-Augmented Generation (RAG) paradigm with the `gemini-2.5-flash` Large Language Model (LLM), the platform autonomously deploys an agentic entity known as the "Yaksha Bot." This agent intercepts community questions and instantly generates highly accurate, context-bound answers derived strictly from the verified FAQ database.

To guarantee platform integrity, the system integrates a synchronous **AI Toxicity Moderation** layer that intercepts and blocks abusive content pre-database insertion. This is complemented by a gamified **Expert Promotion & Reputation Engine** that mathematically rewards users for contributing high-signal answers. Finally, to achieve parity with modern single-page applications (SPAs), all system interactions are broadcast instantaneously via **Socket.io WebSockets**, establishing a seamless, event-driven user experience.

2 Introduction

2.1 Background and Motivation

In large-scale environments, user inquiries typically follow a Pareto distribution — approximately 80% of support interactions originate from a recurring 20% of common procedural questions. While static FAQ pages attempt to mitigate this operational overhead, they fail to account for the natural variance in human morphology and syntax. A user inquiring, "How do I acquire my clearance document?" will inherently fail to trigger a keyword match for an official document titled "No Objection Certificate Protocol," resulting in search abandonment.

Conversely, while community-driven forums (such as StackOverflow, Reddit, or Quora) successfully decentralize knowledge sharing by allowing peers to answer questions, they introduce three critical failure vectors:

1. **Response Latency:** Inquiries rely on human availability. Users may wait hours or days for an expert to view and answer their post.
2. **Content Toxicity:** Open platforms invite bad actors. Unmoderated forums rapidly degenerate into spam repositories without expensive manual oversight.
3. **Information Entropy:** As multiple users provide conflicting answers, the absolute ground truth becomes obscured, requiring users to guess which peer is providing accurate policy data.

2.2 Proposed Solution: Samagama

Samagama (Sanskrit for "coming together" or "confluence") was engineered to explicitly solve these deficiencies by bridging authoritative institutional knowledge with open community collaboration. By unifying the MERN stack (MongoDB, Express, React/Next.js, Node.js) with Google's Generative AI APIs, the project establishes a self-sustaining, low-latency knowledge loop.

It provides an authoritative, semantically searchable FAQ board managed exclusively by administrators, seamlessly integrated with an open Community Q&A hub. The strategic integration of LLMs shifts the burden of moderation and initial query resolution from human administrators to deterministic algorithms. This drastically reduces operational expenditures while providing users with instantaneous, accurate, and cryptographically secure information retrieval.

2.3 System Requirements

The following table outlines the minimum hardware and software prerequisites for deploying and running the platform.


Requirement Type	Specification
Frontend Framework	Next.js 14+ (App Router), React 18, Tailwind CSS
Backend Runtime	Node.js v18.x+, Express.js
Database System	MongoDB 6.0+ (Mongoose ODM)
AI Integration	Google Generative AI SDK (Gemini 2.5 Flash, text-embedding-004)
Real-time Comm.	Socket.io (WebSocket Protocol)
Authentication	Passport.js (Google OAuth 2.0), JWT (JSON Web Tokens)

Table 1: Software and Technology Stack Requirements

3 System Design & Architecture


3.1 High-Level Architecture

The system architecture adheres to a strict Client-Server separation, augmented by an external Machine Learning Intelligence layer. The Next.js frontend handles Server-Side Rendering (SSR) for Search Engine Optimization (SEO) parity, while the Express.js backend acts as the central orchestrator for asynchronous database operations, WebSocket multiplexing, and AI API consumption.




3.2 Entity Relationship Diagram (ERD)

The database consists of intricately linked entities, specifically engineered to support gamification (reputation tracking) and real-time nested commentary.



3.3 Use-Case Diagram

The application services two primary human actors: Citizens (end-users) and Administrators, supplemented by the Yaksha AI Bot acting as an autonomous internal actor responsible for algorithmic enforcement and generation.




3.3.1 Detailed Use Case Specification: RAG Auto-Answer Generation

To provide deeper architectural context, the critical AI Auto-Answer use case is formally specified below:

Use Case ID & Name	UC-08: RAG Auto-Answer Generation
Primary Actor	Yaksha AI Bot (Autonomous System)
Secondary Actors	Google Gemini API, Active Socket.io Clients
Preconditions	1. User successfully posts a question. 2. Content passes Synchronous Toxicity Screening.
Main Success Scenario	1. System intercepts the new question event. 2. System calls Gemini API to convert question text into a 768-D vector. 3. MongoDB executes a <code>\$vectorSearch</code> cosine similarity pipeline against the FAQ collection. 4. System retrieves the top 3 semantically matching official FAQs. 5. System injects the FAQs as context into a prompt payload. 6. Gemini 2.5 Flash API generates a synthesized Markdown response. 7. System saves the Answer document to MongoDB marked as <code>isAI: true</code>. 8. Socket.io broadcasts the new answer to all users actively viewing the thread.
Exceptions (Alternate)	- API Timeout: If Gemini API takes $> 3000\text{ms}$, the asynchronous job fails gracefully and logs an error; the thread remains open for human answers. - Low Cosine Score: If the highest matching FAQ vector score is < 0.65, the system aborts generation to prevent hallucinated answers.
Postconditions	A verified AI answer is permanently attached to the Question Thread, accelerating issue resolution.


3.4 Machine Learning Inference Pipeline

The RAG (Retrieval-Augmented Generation) inference loop is central to the platform's autonomous capabilities, allowing the system to instantly intercept questions and synthesize answers using vector similarity.




3.5 API Route & Middleware Design

The Node.js backend utilizes a strict horizontal routing pipeline. Every incoming HTTP request is funneled through security and validation middleware before reaching the core domain controllers.



3.6 End-to-End System Pipeline Flow

The complete macroscopic pipeline demonstrates how traditional CRUD operations synchronize with the ML Inference loop and real-time Socket.io state management.




4 Data Flow Diagrams (DFD)

To comprehensively map the movement of data across the Samagama platform, we utilize Data Flow Diagrams at multiple levels of abstraction.


4.1 Level 0 DFD (Context Diagram)

The Level 0 DFD provides a macroscopic view of the entire system, illustrating how external entities interact with the core Samagama platform.



4.2 Level 1 DFD

The Level 1 DFD decomposes the monolithic system into its primary subsystems: Authentication, Knowledge Base Management, and the AI Inference Pipeline.



5 Implementation

5.1 Tech Stack & Justification

The selection of the MERN stack coupled with Google's Generative AI suite was a highly calculated architectural decision optimized for throughput, developmental velocity, and geometric scalability.

- **Frontend Framework (Next.js & React 18):** Next.js was selected specifically for its hybrid rendering capabilities. Because public FAQ documents and Community Q&A threads must be rapidly indexable by Google and Bing search crawlers, Server-Side Rendering (SSR) is strictly necessary. Traditional Create-React-App SPAs return blank HTML payloads, crippling SEO. Furthermore, Tailwind CSS allows for rapid, constraint-based styling via utility classes, crucial for implementing the complex glassmorphic aesthetics and responsive grid layouts without CSS bloat.
- **Backend Environment (Node.js & Express.js):** Node.js utilizes an event-driven, non-blocking I/O model based on the V8 engine. This architecture is critical for managing thousands of persistent WebSocket connections concurrently alongside intensive HTTP REST traffic. A multithreaded environment (like Java Spring) would suffer from massive memory overhead per socket connection, whereas Node handles WebSockets natively with minimal resource consumption.
- **Database Layer (MongoDB & Mongoose ORM):** The data structures inherent in a Q&A platform are highly hierarchical and nested (e.g., Questions containing arrays of Answers, which contain arrays of Comments and Upvote Metadata). A NoSQL document-oriented database like MongoDB handles deep JSON nesting natively. Most importantly, MongoDB seamlessly supports storing 768-dimensional float arrays (the semantic vectors) directly on the document schema, facilitating high-speed retrieval without requiring a secondary, isolated Vector database architecture for this scale.
- **Machine Learning Intelligence (Google Gemini APIs):** Google provides the `text-embedding-004` model for state-of-the-art semantic dimensionality at incredibly low latency. For text generation, the `gemini2.5flash` model offers exceptional reasoning capabilities with near-instantaneous token generation (often under 2 seconds). This speed is essential for synchronous operations like toxicity screening,

where the HTTP request cannot hang for more than 500ms without degrading the User Experience (UX).

- **Real-Time Telemetry (Socket.io):** Raw WebSockets (ws protocol) lack automatic reconnection fallback mechanisms and complex room multiplexing. Socket.io natively supports "Namespaces" and "Rooms", allowing the Express server to broadcast "User is Typing" events exclusively to the subset of users currently viewing a specific Question Thread, rather than globally flooding the entire network.

5.2 API Design & Endpoints

The system follows strict RESTful principles, utilizing standard HTTP verbs and status codes for predictable resource modification.

Endpoint	Method	Description & Auth Level
/api/auth/google	GET	Initiates the Google OAuth 2.0 flow. (Public)
/api/faqs/search	GET	Executes vector similarity search against the DB. (Public)
/api/qa	POST	Submits a new community question. Invokes Toxicity check. (User)
/api/qa/:id/answers	POST	Posts an answer. Evaluates upvote mechanics. (User)
/api/faqs	POST	Creates a new official FAQ entry and generates embedding. (Admin)
/api/analytics/stats	GET	Fetches platform telemetry (User count, active threads). (Admin)
/api/chatbot	POST	Direct RAG generation for the floating UI Yaksha chatbot. (Public)

Table 2: Core REST API Endpoints

5.3 Module & Feature Breakdown

The implementation logic is partitioned into highly cohesive, loosely coupled micro-modules within the monolith:

1. **Authentication Controller (authController.js):** Manages Google OAuth 2.0 handshakes and local bcryptjs cryptographic password hashing. It issues JSON

Web Tokens (JWT) inside `HttpOnly` cookies to strictly mitigate Cross-Site Scripting (XSS) extraction attacks.

2. **FAQ Controller** (`faqController.js`): Manages complete CRUD operations for the core institutional knowledge base. This module contains the critical pipeline `generateEmbedding()` that vectorizes textual payloads and persists them to the FAQ MongoDB collection.
3. **Q&A Controller** (`qaController.js`): The most complex orchestration module. It handles user submissions, recursively populates author metadata, manages the distributed lock logic for reputation incrementation (upvoting), and invokes the background asynchronous AI Auto-Answer daemon.
4. **Socket Manager** (`server.js`): Wraps the Express `http.Server` instance to establish the `Upgrade` WebSocket connection pool and manage room subscriptions.

6 Feature Spotlight

This section extensively investigates the standout algorithmic innovations of the Samagama platform. It details the theoretical purpose, architectural design, execution logic, mathematical foundations, and real-world impact of each flagship feature.

6.1 Semantic Vector Search (Intent Engine)

6.1.1 Theoretical Purpose

Traditional search architecture utilizes lexical matching algorithms (such as SQL LIKE '%keyword%' or Elasticsearch BM25). These algorithms suffer from a fundamental flaw: they measure string overlap, not meaning. If a student searches for "When will I receive my financial aid?", but the official FAQ document states "Stipends are disbursed on the 1st of every month", a lexical search will return a zero-match because the strings do not geometrically overlap.

Semantic Vector Search solves this by converting human language into highdimensional mathematical geometry, mapping arbitrary strings to a continuous vector space where distance correlates to semantic meaning.

6.1.2 Architecture and Execution Algorithm


When an Administrator creates or updates an FAQ, the system extracts the text, concatenates the Question and Answer, and transmits it via gRPC to the `text-embedding-004` model. Gemini returns a dense vector $V \in \mathbb{R}^{768}$. This 768-dimensional float array is serialized and saved directly to the MongoDB document.

When a user submits a search query Q , the system instantly vectorizes Q into query vector V_Q . The Node.js backend retrieves all FAQ vectors V_{F_i} from the database and executes a high-speed In-Memory Cosine Similarity calculation. The Cosine Similarity equation determines the angle between the two vectors in 768-dimensional space, effectively measuring conceptual closeness regardless of vector magnitude (text length).

$$\text{Cosine Similarity}(V_Q, V_{F_i}) = \cos(\theta) = \frac{V_Q \cdot V_{F_i}}{\|V_Q\| \|V_{F_i}\|} = \frac{\sum_{j=1}^{768} V_{Q_j} V_{F_{i,j}}}{\sqrt{\sum_{j=1}^{768} V_{Q_j}^2} \sqrt{\sum_{j=1}^{768} V_{F_{i,j}}^2}} \quad (1)$$

The algorithm scores every FAQ on a scale from -1.0 (complete opposites) to 1.0 (perfect match). The system applies a rigid empirical threshold, dropping any result scoring

below 0.55 to prevent hallucinated matches, and returns the highest-scoring documents array to the client.



6.1.3 Real-World Impact

Usability Impact: Eliminates the "No Results Found" dead-end. Users locate the exact institutional policy they require regardless of how poorly, abstractly, or colloquially they formulate their search query. This drastically reduces the volume of duplicate questions posted to the community hub, saving hundreds of hours of administrative triaging.


6.2 Deep Auto-Answer Pipeline (Yaksha Bot)

6.2.1 Theoretical Purpose

In decentralized community forums, users frequently ask highly specific, repetitive questions that are technically covered by existing FAQs, but they bypass the search bar entirely out of convenience. Instead of waiting hours for a human moderator to manually locate the FAQ, copy the link, and paste it as a reply, the platform deploys an autonomous AI agent (Yaksha). This agent instantly reads the user's question, locates the relevant policy, and synthesizes a custom, polite response.

6.2.2 Sequence Diagram & RAG Flow

This diagram outlines the precise chronological flow of operations when a user submits a question and the system autonomously generates a response via Retrieval-Augmented Generation.



6.2.3 Architecture and RAG Algorithm

To deploy Generative AI in an institutional setting without the catastrophic risk of the LLM "hallucinating" false policies or providing generic web data, this feature imple-

ments strict **Retrieval-Augmented Generation (RAG)**. The LLM is algorithmically forbidden from relying on its internal parametric memory weights.

1. **Trigger Event:** A user submits a new question to the Community Q&A. The HTTP POST request resolves immediately (201 `Created`) to prevent blocking the client UI thread.
2. **Retrieval Phase:** A background daemon vectorizes the new question and performs a semantic search against the authoritative FAQ database. It retrieves the top 3 highly relevant FAQs. If no FAQs score above an aggressive 0.65 threshold, the daemon aborts execution (the Yaksha Bot remains silent if it lacks the factual data to answer).
3. **Augmentation Phase:** The daemon extracts the plaintext of the retrieved FAQs and concatenates them into a rigid, structured `CONTEXT` payload block.
4. **Generation Phase:** `gemini-2.5-flash` is invoked with a strict system instruction: *"Act as the Yaksha Bot. Answer the user's question using EXCLUSIVELY the provided CONTEXT. If the context does not explicitly contain the answer, reply that you do not know."*
5. **Delivery Phase:** The synthesized Markdown answer is persisted to MongoDB with an `isAI: true` Boolean flag. A Socket.io event triggers the frontend to dynamically render the response, augmenting it with a specialized UI badge to indicate machine origin.

6.2.4 Real-World Impact

Resolution Impact: Eradicates response latency entirely. A student inquiring about a document procedure at 3:00 AM receives a perfectly formatted, polite, and procedurally accurate response from the Yaksha Bot within 3.5 seconds. This closes the support loop instantaneously without requiring human administrative intervention.

Algorithm 1 Deep Auto-Answer Daemon (RAG Implementation)

```

1: procedure YAKSHAAUTOANSWER(question_id, title, body)
2:   query_vector  $\leftarrow$  GENERATEEMBEDDING(title + body)
3:   all_faqs  $\leftarrow$  FETCHALLPUBLISHEDFAQS
4:   similarity_scores  $\leftarrow$  COSINESIMILARITY(query_vector, all_faqs)
5:   top_matches  $\leftarrow$  FILTER(similarity_scores, score  $\geq$  0.65)
6:   if LENGTH(top_matches) > 0 then
7:     context_block  $\leftarrow$  CONCATENATETEXT(top_matches)
8:     prompt  $\leftarrow$  "System Directive: Answer using ONLY the CONTEXT."
9:     prompt  $\leftarrow$  prompt + "\nContext: " + context_block
10:    prompt  $\leftarrow$  prompt + "\nUser Question: " + title
11:    ai_response  $\leftarrow$  INVOKEGEMINIFLASH(prompt)
12:    PERSISTANSWERTODB(question_id, ai_response, isAI = true)
13:    BROADCASTSOCKETEVENT("new_answer_notification")
14:  else
15:    return ▷ Abort silently. Do not hallucinate.
16:  end if
17: end procedure

```

6.3 Synchronous AI Toxicity Moderation


6.3.1 Theoretical Purpose

Open crowdsourced platforms are notoriously vulnerable to bad actors distributing abusive, sexually explicit, highly toxic, or completely irrelevant spam content. Scaling a human moderation team to monitor a 24/7 forum is financially unviable. Furthermore, retroactive moderation (deleting toxic posts after they are published) allows the damage to occur. By placing a synchronous AI gate in front of the database ingestion point, the system acts as an impenetrable, proactive shield.

6.3.2 Architecture and Execution Algorithm

Unlike the Auto-Answer pipeline (which is executed asynchronously), Toxicity Screening must be **synchronous** (blocking). If a user attempts to persist toxic content to the database, the HTTP request must fail immediately during the transaction.

To achieve this without causing request timeouts, the system utilizes the `gemini-2.5flash` model due to its extreme Time-to-First-Token (TTFT) low latency. The text payload (Title + Body) is injected into the model with the zero-shot prompt: *"Analyze this text. Is it abusive, highly toxic, or sexual? Reply strictly with exactly SAFE or TOXIC."* If the model returns `TOXIC`, the Node.js controller halts execution and immediately throws a `400 Bad Request`.



6.3.3 Real-World Impact

Security Impact: Guarantees a pristine, professional digital environment. The platform remains safe for all users and minors without requiring a dedicated 24/7 human moderation team, simultaneously protecting institutional liability and brand integrity.

6.4 Expert Promotion & Reputation Engine

6.4.1 Theoretical Purpose

To foster a thriving, self-sustaining community, domain experts must be psychologically and socially incentivized to spend their personal time answering questions. The Reputation Engine introduces gamification mechanics to knowledge sharing. Mirroring successful platforms like StackOverflow, users accrue points (REP) for providing helpful, accurate answers, transforming community participation into a visible status symbol.

6.4.2 Architecture and Execution Algorithm

The `User` schema in MongoDB tracks an integer `reputation` score.

- **Upvoting Mechanics (+10 REP):** Users can mathematically endorse helpful questions or answers. To prevent system abuse (e.g., a user clicking upvote 100 times), the system implements a strict **idempotency guard**. The database schema maintains an array of `upvotedBy` ObjectIDs. If a user attempts to upvote twice, the backend controller ignores the request. If the user toggles their upvote off, the array pulls their ID and atomic operations mathematically subtract 10 REP from the target author.
- **Accepted Answer Mechanics (+20 REP):** The original author of a Question possesses the exclusive privilege to click a "Mark as Accepted" Boolean switch on the highest quality answer. This sets an `isAccepted: true` flag on the Answer document. The frontend renders a permanent green border around the answer, and the backend permanently transfers an additional +20 REP to the answering user, visually distinguishing the absolute ground truth for future readers.

6.4.3 Real-World Impact

Engagement Impact: Drives massive user retention and engagement. By gamifying accuracy, the system naturally floats the most accurate information to the top of the thread hierarchy via democratic consensus. As experts accumulate REP, they gain verifiable community trust.

6.5 Live Socket.io Real-Time Telemetry

6.5.1 Theoretical Purpose


Modern web applications must feel alive. If User A is typing an answer to User B's question, User B should not be required to manually refresh the browser page to discover if a response has arrived. Real-time telemetry maintains deep user immersion and dramatically reduces bounce rates.

6.5.2 Architecture and Execution Algorithm

The Express.js server integrates `socket.io`. Upon initial client connection, the server maps the WebSocket instance to the user's authenticated ID via `socket.join('user_' + userId)`, establishing a private communication channel.

When users navigate to a specific Q&A thread URL (e.g., `/qa/123`), the React `useEffect` hook emits a join event for a specific multiplexed "Room" (`socket.join('thread_123')`).

When a user triggers an `onChange` event in the text area, a `typing` payload is emitted to the server. The server instantly broadcasts a `user_typing` event *exclusively* to the `thread_123` room. When the answer is finally submitted, the server saves it to MongoDB and emits a `new_notification` event to the original author's private `user_X` room, triggering a UI toast notification.



6.5.3 Real-World Impact

UX Impact: Replicates the fluid, highly interactive experience of modern enterprise platforms (like Slack or Microsoft Teams). Users receive instant visual feedback (toast notifications, animated typing bubbles), drastically increasing time-on-site and aggregate platform satisfaction.

7 Challenges & System Limitations

7.1 Challenge 1: AI Latency vs. UX Blocking

Problem: Generating a high-quality Markdown answer via the Gemini LLM requires approximately 3 to 4 seconds of network and processing time. If the Auto-Answer pipeline were executed synchronously during the HTTP POST request, the user's browser would freeze for 4 seconds. This leads to UX degradation, prompting users to repeatedly click the "Submit" button, resulting in duplicated database entries and corrupted state.

Architectural Solution: The architecture entirely decoupled the AI generation phase from the HTTP request phase. The backend saves the initial question to MongoDB and returns a 201 `Created` response within 50ms. The AI RAG pipeline is then executed as a detached asynchronous background promise. Once completed, the daemon utilizes the Socket.io connection to push the finalized answer to the client seamlessly.

7.2 Challenge 2: Vector Dimensionality Mismatch

Problem: During initial prototyping, experimenting with different embedding models caused fatal mathematical crashes in the Cosine Similarity function due to mismatched array lengths (e.g., attempting to execute a dot-product between a 768-dimensional array and a 1536-dimensional array).

Architectural Solution: Strict schema validation was implemented at the Mongoose ORM layer to assert that the `embedding` field is strictly an Array of Numbers. Furthermore, the system enforces exclusive, hardcoded use of `text-embedding-004` across all ingestion and retrieval points to guarantee geometrical parity.

7.3 Limitation: In-Memory Array Iteration

System Limitation: Currently, the Semantic Vector Search fetches all FAQ vectors from MongoDB into Node.js V8 RAM and calculates Cosine Similarity using standard JavaScript mapping. While this executes in under 10ms for hundreds of FAQs, it fundamentally operates at $O(N)$ time complexity. If the knowledge base scales to 1,000,000+ FAQs, calculating one million dot-products sequentially in JavaScript will catastrophically bottleneck the single-threaded Node Event Loop.

8 Future Architectural Enhancements

- **Dedicated Vector Database Indexing (Atlas Vector Search):** To mathematically solve the $O(N)$ CPU limitation discussed above, the system architecture can be upgraded to utilize MongoDB Atlas Vector Search natively. This utilizes Hierarchical Navigable Small Worlds (HNSW) algorithms to find Approximate Nearest Neighbors (ANN) in $O(\log N)$ time complexity, allowing the platform to scale to millions of high-dimensional documents effortlessly without degrading search latency.
- **Multilingual Embedding Pipelines:** The global user base is highly linguistically diverse. By hot-swapping the embedding model to a multilingual variant (e.g., `text-multilingual-embedding-002`), a user could execute a search query in Hindi or Spanish, and the Intent Engine would successfully retrieve the semantically matching English FAQ, completely breaking down language barriers in institutional support.
- **Generative UI & Deterministic Citations:** The Yaksha Bot currently synthesizes pure Markdown text. In future iterations, it could be upgraded to generate interactive React Server Components (Generative UI), allowing it to embed clickable, routed citations back to the specific authoritative FAQ documents it utilized as context in its RAG pipeline. This increases transparency and auditability.
- **Multimodal Image Support / Optical Character Recognition (OCR):** Allowing users to upload screenshots of procedural errors. The backend could pass the image array buffer to `gemini-1.5-pro-vision` to extract the context, read the error code, and autonomously answer visual bugs that users struggle to describe via text.

9 Conclusion

The Samagama FAQ & Crowdsourcing Platform successfully and completely modernizes the concept of institutional knowledge management. By transitioning from a static, rigid, keyword-reliant webpage into a highly dynamic, AI-shielded ecosystem, the project decisively solves the critical issues of search failure and delayed community resolution that plague legacy systems.

The integration of **Semantic Vector Search** ensures that users locate the procedural information they require based on abstract meaning, rather than exact lexical phrasing. The **Deep Auto-Answer Pipeline** transforms the system from a passive ledger into an active, intelligent assistant that resolves complex queries instantly via the "Yaksha Bot", while completely neutralizing the threat of LLM hallucinations through strict Retrieval-Augmented Generation (RAG) constraints.

Furthermore, by combining these advanced Generative AI capabilities with highly interactive, real-time WebSockets and gamified reputation engines characteristic of the modern MERN stack, Samagama delivers an exceptional, scalable, and cryptographically safe user experience. It stands as a comprehensive, production-ready blueprint for how modern institutions can leverage machine learning not merely for offline data analysis, but for direct, real-time user empowerment and operational efficiency.